

Heart Rate Modelling

December 16, 2021

Model building to predict heart disease using logistic Regression

```
[1]: #Import modules
import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
from sklearn.metrics import log_loss
from sklearn.metrics import jaccard_score
from scipy import stats
%matplotlib inline
import matplotlib.pyplot as plt
```

```
[2]: #Load Data
heart_df=pd.read_csv(r'C:\Users\sanus\Documents\Programming\Python\Python data_\
↪analysis portfolio\Data\heart.csv')
heart_df.head()
```

```
[2]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

```
[3]: #Explore Data types
heart_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   918 non-null   int64
1   Sex                   918 non-null   object
2   ChestPainType         918 non-null   object
3   RestingBP             918 non-null   int64
4   Cholesterol            918 non-null   int64
5   FastingBS             918 non-null   int64
6   RestingECG            918 non-null   object
7   MaxHR                 918 non-null   int64
8   ExerciseAngina        918 non-null   object
9   Oldpeak               918 non-null   float64
10  ST_Slope              918 non-null   object
11  HeartDisease          918 non-null   int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
[4]: #optional: seperate categorical variables from numerical variables
#def get_numerical_var(dframe):
#     lis=[]
#     for col in heart_df:
#         if heart_df[col].dtype!=object:
#             header='' + col
#             lis.append(header)
#     return lis

#def get_cat_var(dframe):
#     lis=[]
#     for col in heart_df:
#         if heart_df[col].dtype==object:
#             header='' + col
#             lis.append(header)
#     return lis

#print(get_cat_var(heart_df))
```

```
[5]: #Convert categorical independent variables to dummy variables
le = LabelEncoder()
heart_df['Sex']=le.fit_transform(heart_df['Sex'])
heart_df['ChestPainType']=le.fit_transform(heart_df['ChestPainType'])
heart_df['RestingECG']=le.fit_transform(heart_df['RestingECG'])
heart_df['ExerciseAngina']=le.fit_transform(heart_df['ExerciseAngina'])
```

```
heart_df['ST_Slope']=le.fit_transform(heart_df['ST_Slope'])
```

```
[6]: #Visualize using boxplots
# fig=plt.figure()
# ax0=fig.add_subplot(1,2,1)
# heart_df.plot(kind='box',
#                 figsize=(30,10),
#                 ax=ax0)
# ax0.set_title('Box plots ')
# plt.show
```

```
[7]: #Calculate pearson coefficient-optional
heart_df.corr()
```

```
[7]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	\
Age	1.000000	0.055750	-0.077150	0.254399	-0.095282	
Sex	0.055750	1.000000	-0.126559	0.005133	-0.200092	
ChestPainType	-0.077150	-0.126559	1.000000	-0.020647	0.067880	
RestingBP	0.254399	0.005133	-0.020647	1.000000	0.100893	
Cholesterol	-0.095282	-0.200092	0.067880	0.100893	1.000000	
FastingBS	0.198039	0.120076	-0.073151	0.070193	-0.260974	
RestingECG	-0.007484	0.071552	-0.072537	0.022656	-0.196544	
MaxHR	-0.382045	-0.189186	0.289123	-0.112135	0.235792	
ExerciseAngina	0.215793	0.190664	-0.354727	0.155101	-0.034166	
Oldpeak	0.258612	0.105734	-0.177377	0.164803	0.050148	
ST_Slope	-0.268264	-0.150693	0.213521	-0.075162	0.111471	
HeartDisease	0.282039	0.305445	-0.386828	0.107589	-0.232741	

	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	\
Age	0.198039	-0.007484	-0.382045	0.215793	0.258612	
Sex	0.120076	0.071552	-0.189186	0.190664	0.105734	
ChestPainType	-0.073151	-0.072537	0.289123	-0.354727	-0.177377	
RestingBP	0.070193	0.022656	-0.112135	0.155101	0.164803	
Cholesterol	-0.260974	-0.196544	0.235792	-0.034166	0.050148	
FastingBS	1.000000	0.087050	-0.131438	0.060451	0.052698	
RestingECG	0.087050	1.000000	-0.179276	0.077500	-0.020438	
MaxHR	-0.131438	-0.179276	1.000000	-0.370425	-0.160691	
ExerciseAngina	0.060451	0.077500	-0.370425	1.000000	0.408752	
Oldpeak	0.052698	-0.020438	-0.160691	0.408752	1.000000	
ST_Slope	-0.175774	-0.006778	0.343419	-0.428706	-0.501921	
HeartDisease	0.267291	0.057384	-0.400421	0.494282	0.403951	

	ST_Slope	HeartDisease
Age	-0.268264	0.282039
Sex	-0.150693	0.305445
ChestPainType	0.213521	-0.386828
RestingBP	-0.075162	0.107589

Cholesterol	0.111471	-0.232741
FastingBS	-0.175774	0.267291
RestingECG	-0.006778	0.057384
MaxHR	0.343419	-0.400421
ExerciseAngina	-0.428706	0.494282
Oldpeak	-0.501921	0.403951
ST_Slope	1.000000	-0.558771
HeartDisease	-0.558771	1.000000

```
[8]: #Calculate p-value
for col in heart_df.columns:
    pearson_coef, p_value = stats.pearsonr(heart_df[col],
    ↪heart_df['HeartDisease'])
    print("For", col, "The Pearson Correlation Coefficient is", pearson_coef, "
    ↪with a P-value of P =", p_value)
```

For Age The Pearson Correlation Coefficient is 0.28203850581899736 with a P-value of P = 3.007953240047123e-18

For Sex The Pearson Correlation Coefficient is 0.30544491596314066 with a P-value of P = 2.821897823681047e-21

For ChestPainType The Pearson Correlation Coefficient is -0.38682769426256447 with a P-value of P = 3.8887950200777145e-34

For RestingBP The Pearson Correlation Coefficient is 0.10758898037140399 with a P-value of P = 0.0010953145851714478

For Cholesterol The Pearson Correlation Coefficient is -0.2327406389270114 with a P-value of P = 9.308308883525767e-13

For FastingBS The Pearson Correlation Coefficient is 0.267291186110298 with a P-value of P = 1.7535980103286795e-16

For RestingECG The Pearson Correlation Coefficient is 0.05738435701345111 with a P-value of P = 0.08225947154724081

For MaxHR The Pearson Correlation Coefficient is -0.4004207694631902 with a P-value of P = 1.137785984026953e-36

For ExerciseAngina The Pearson Correlation Coefficient is 0.49428199182426846 with a P-value of P = 1.0130182683908042e-57

For Oldpeak The Pearson Correlation Coefficient is 0.4039507220628864 with a P-value of P = 2.3907724240568936e-37

For ST_Slope The Pearson Correlation Coefficient is -0.5587707148497059 with a P-value of P = 1.6715991289946643e-76

For HeartDisease The Pearson Correlation Coefficient is 0.9999999999999998 with a P-value of P = 0.0

```
[9]: for col in heart_df.columns:
    strong_corr=''
    if p_value <0.001:
```

```
strong_corr + ','+ col
print(strong_corr)
```

```
[10]: #The p_value for all variables show their statistical relationship is
      ↪significant
```

```
[11]: #Define X and convert to numpy array
X = np.asarray(heart_df[['Age', 'Sex', 'ChestPainType', 'RestingBP',
      ↪'Cholesterol', 'FastingBS',
      ↪'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak', 'ST_Slope']])
X[0:5]
```

```
[11]: array([[ 40. ,   1. ,   1. , 140. , 289. ,   0. ,   1. , 172. ,   0. ,
           0. ,   2. ],
          [ 49. ,   0. ,   2. , 160. , 180. ,   0. ,   1. , 156. ,   0. ,
           1. ,   1. ],
          [ 37. ,   1. ,   1. , 130. , 283. ,   0. ,   2. ,   98. ,   0. ,
           0. ,   2. ],
          [ 48. ,   0. ,   0. , 138. , 214. ,   0. ,   1. , 108. ,   1. ,
           1.5,   1. ],
          [ 54. ,   1. ,   2. , 150. , 195. ,   0. ,   1. , 122. ,   0. ,
           0. ,   2. ]])
```

```
[12]: #Define Y and convert to numpy array
y = np.asarray(heart_df['HeartDisease'])
y [0:5]
```

```
[12]: array([0, 1, 0, 1, 0], dtype=int64)
```

```
[13]: #Normalize
X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

```
[13]: array([[ -1.4331398 ,   0.51595242,   0.22903206,   0.41090889,   0.82507026,
          -0.55134134,   0.01725451,   1.38292822,  -0.8235563 ,  -0.83243239,
```

```

1.05211381],
[-0.47848359, -1.93816322, 1.27505906, 1.49175234, -0.17196105,
-0.55134134, 0.01725451, 0.75415714, -0.8235563 , 0.10566353,
-0.59607813],
[-1.75135854, 0.51595242, 0.22903206, -0.12951283, 0.7701878 ,
-0.55134134, 1.60121899, -1.52513802, -0.8235563 , -0.83243239,
1.05211381],
[-0.5845565 , -1.93816322, -0.81699495, 0.30282455, 0.13903954,
-0.55134134, 0.01725451, -1.13215609, 1.21424608, 0.57471149,
-0.59607813],
[ 0.05188098, 0.51595242, 1.27505906, 0.95133062, -0.0347549 ,
-0.55134134, 0.01725451, -0.5819814 , -0.8235563 , -0.83243239,
1.05211381]])

```

Train and Predict

```

[14]: #Train Model
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2,
↳random_state=4)
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)

```

Train set: (734, 11) (734,)

Test set: (184, 11) (184,)

```

[15]: LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
LR

```

```

[15]: LogisticRegression(C=0.01, solver='liblinear')

```

```

[16]: #Predict
yhat = LR.predict(X_test)
yhat

```

```

[16]: array([0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,
0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1,
1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0,
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1,
1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0,
1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1,
1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0,
0, 1, 1, 1, 1, 0, 1, 1], dtype=int64)

```

```

[17]: yhat_prob = LR.predict_proba(X_test)
yhat_prob

```

```

[17]: array([[0.5156815 , 0.4843185 ],
[0.21000747, 0.78999253],

```

[0.56400433, 0.43599567],
[0.75538125, 0.24461875],
[0.67403947, 0.32596053],
[0.18052657, 0.81947343],
[0.21522349, 0.78477651],
[0.38227614, 0.61772386],
[0.54362217, 0.45637783],
[0.67546032, 0.32453968],
[0.89131343, 0.10868657],
[0.51782803, 0.48217197],
[0.80825933, 0.19174067],
[0.77862184, 0.22137816],
[0.64098133, 0.35901867],
[0.07528364, 0.92471636],
[0.84393956, 0.15606044],
[0.79067292, 0.20932708],
[0.83189162, 0.16810838],
[0.25299, 0.74701],
[0.19278601, 0.80721399],
[0.71229155, 0.28770845],
[0.80218114, 0.19781886],
[0.78344398, 0.21655602],
[0.08958179, 0.91041821],
[0.26834287, 0.73165713],
[0.89460813, 0.10539187],
[0.50407228, 0.49592772],
[0.20848928, 0.79151072],
[0.54787334, 0.45212666],
[0.1516258, 0.8483742],
[0.24292488, 0.75707512],
[0.05299495, 0.94700505],
[0.21291484, 0.78708516],
[0.44016292, 0.55983708],
[0.78987214, 0.21012786],
[0.15123149, 0.84876851],
[0.04791338, 0.95208662],
[0.10859109, 0.89140891],
[0.60106079, 0.39893921],
[0.11901581, 0.88098419],
[0.84462885, 0.15537115],
[0.32267183, 0.67732817],
[0.40499055, 0.59500945],
[0.25152373, 0.74847627],
[0.17098681, 0.82901319],
[0.20328074, 0.79671926],
[0.19043815, 0.80956185],
[0.66150182, 0.33849818],

[0.66226393, 0.33773607],
[0.40262852, 0.59737148],
[0.66933417, 0.33066583],
[0.8348269 , 0.1651731],
[0.88946666, 0.11053334],
[0.86068682, 0.13931318],
[0.68396416, 0.31603584],
[0.10240867, 0.89759133],
[0.53492133, 0.46507867],
[0.85390378, 0.14609622],
[0.14393647, 0.85606353],
[0.29518045, 0.70481955],
[0.89662348, 0.10337652],
[0.82326528, 0.17673472],
[0.46494034, 0.53505966],
[0.05511084, 0.94488916],
[0.13321716, 0.86678284],
[0.50561048, 0.49438952],
[0.16822714, 0.83177286],
[0.85267776, 0.14732224],
[0.88173901, 0.11826099],
[0.18824383, 0.81175617],
[0.50137339, 0.49862661],
[0.82584709, 0.17415291],
[0.7507718 , 0.2492282],
[0.75475291, 0.24524709],
[0.56378538, 0.43621462],
[0.67369807, 0.32630193],
[0.71952348, 0.28047652],
[0.78445745, 0.21554255],
[0.16807562, 0.83192438],
[0.60481667, 0.39518333],
[0.22984594, 0.77015406],
[0.64659397, 0.35340603],
[0.27594705, 0.72405295],
[0.02430098, 0.97569902],
[0.77855616, 0.22144384],
[0.82007122, 0.17992878],
[0.86843899, 0.13156101],
[0.7418594 , 0.2581406],
[0.12691496, 0.87308504],
[0.09177076, 0.90822924],
[0.49643833, 0.50356167],
[0.77766977, 0.22233023],
[0.62953108, 0.37046892],
[0.44925408, 0.55074592],
[0.2371203 , 0.7628797],

[0.63163649, 0.36836351],
[0.22654474, 0.77345526],
[0.20027087, 0.79972913],
[0.40354267, 0.59645733],
[0.83328723, 0.16671277],
[0.82784663, 0.17215337],
[0.07569148, 0.92430852],
[0.04257631, 0.95742369],
[0.14131135, 0.85868865],
[0.41661621, 0.58338379],
[0.94640812, 0.05359188],
[0.68222178, 0.31777822],
[0.35717521, 0.64282479],
[0.21316846, 0.78683154],
[0.45110218, 0.54889782],
[0.69151796, 0.30848204],
[0.37879249, 0.62120751],
[0.81933271, 0.18066729],
[0.81976397, 0.18023603],
[0.28519955, 0.71480045],
[0.13852824, 0.86147176],
[0.25483103, 0.74516897],
[0.83169544, 0.16830456],
[0.27601039, 0.72398961],
[0.82124094, 0.17875906],
[0.21147787, 0.78852213],
[0.09115008, 0.90884992],
[0.80103946, 0.19896054],
[0.07444636, 0.92555364],
[0.92777234, 0.07222766],
[0.55824038, 0.44175962],
[0.43736313, 0.56263687],
[0.76361355, 0.23638645],
[0.34972024, 0.65027976],
[0.67635119, 0.32364881],
[0.74389765, 0.25610235],
[0.27165166, 0.72834834],
[0.173697, 0.826303],
[0.43329233, 0.56670767],
[0.37085734, 0.62914266],
[0.36404874, 0.63595126],
[0.23263862, 0.76736138],
[0.41459585, 0.58540415],
[0.11978213, 0.88021787],
[0.60370614, 0.39629386],
[0.14005645, 0.85994355],
[0.83727335, 0.16272665],

```
[0.77506364, 0.22493636],
[0.08960527, 0.91039473],
[0.25834229, 0.74165771],
[0.19470285, 0.80529715],
[0.73340696, 0.26659304],
[0.1442758 , 0.8557242 ],
[0.4660422 , 0.5339578 ],
[0.25845861, 0.74154139],
[0.43792758, 0.56207242],
[0.81870215, 0.18129785],
[0.34347672, 0.65652328],
[0.09296722, 0.90703278],
[0.68809507, 0.31190493],
[0.20150081, 0.79849919],
[0.49617001, 0.50382999],
[0.7304135 , 0.2695865 ],
[0.0974097 , 0.9025903 ],
[0.81323138, 0.18676862],
[0.09195634, 0.90804366],
[0.25791424, 0.74208576],
[0.83826492, 0.16173508],
[0.21704828, 0.78295172],
[0.32954349, 0.67045651],
[0.21515663, 0.78484337],
[0.12536697, 0.87463303],
[0.43843534, 0.56156466],
[0.75150608, 0.24849392],
[0.35584798, 0.64415202],
[0.85524862, 0.14475138],
[0.48364278, 0.51635722],
[0.42790088, 0.57209912],
[0.23741146, 0.76258854],
[0.60601205, 0.39398795],
[0.87690889, 0.12309111],
[0.08098681, 0.91901319],
[0.38488163, 0.61511837],
[0.33923496, 0.66076504],
[0.40438148, 0.59561852],
[0.7826983 , 0.2173017 ],
[0.44250076, 0.55749924],
[0.39789302, 0.60210698]])
```

0.0.1 Evaluation

```
[18]: jaccard_score(y_test, yhat, pos_label=0)
```

```
[18]: 0.7628865979381443
```

```
[19]: import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
print(confusion_matrix(y_test, yhat, labels=[1,0]))
```

```
[[87  9]
 [14 74]]
```

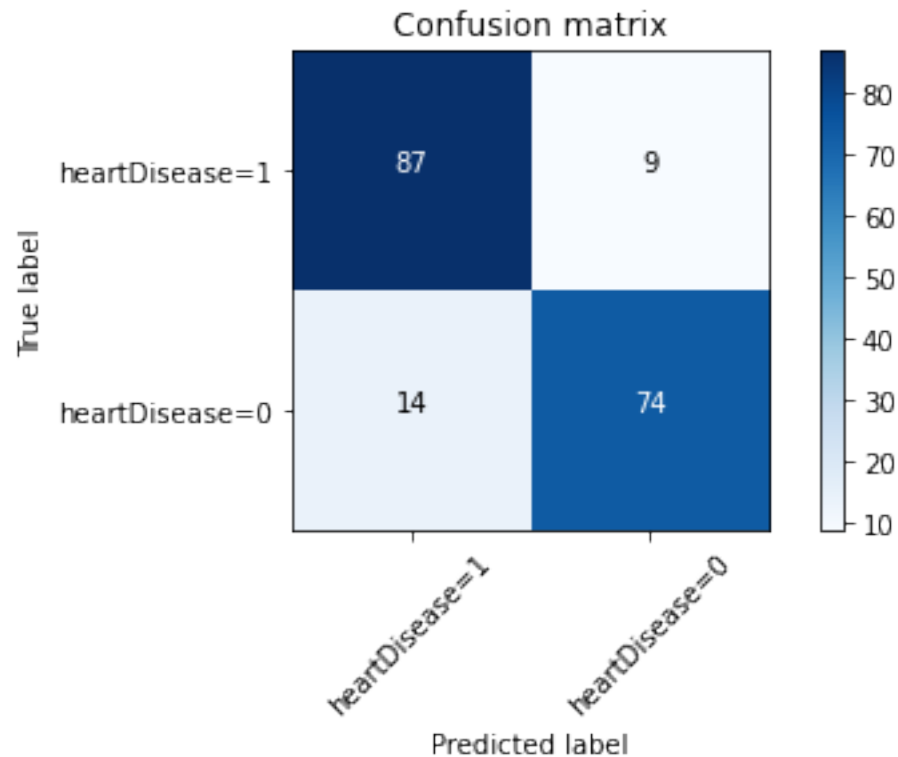
```
[24]: # Compute confusion matrix
cnf_matrix = confusion_matrix(y_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
```

```
plot_confusion_matrix(cnf_matrix,
↳classes=['heartDisease=1','heartDisease=0'],normalize= False,
↳title='Confusion matrix')
```

Confusion matrix, without normalization

```
[[87  9]
 [14 74]]
```



```
[21]: print (classification_report(y_test, yhat))
```

	precision	recall	f1-score	support
0	0.89	0.84	0.87	88
1	0.86	0.91	0.88	96
accuracy			0.88	184
macro avg	0.88	0.87	0.87	184
weighted avg	0.88	0.88	0.87	184

```
[22]: #Calculate log_loss
log_loss(y_test, yhat_prob)
```

[22] : 0.38787602959104506

[]: